

# Based on this business Logic, write the database Schema and use prisma ORM. Determine if we will need a relational database as well, and the choice for this will be Postgres

Of course. Based on the established business logic for the Enterprise Agentic Workflow Engine (EAWE), here is the database schema design using Prisma ORM.

## Database Strategy: A Hybrid Approach

To meet all requirements of the business logic, we will use a hybrid database strategy. A simple relational database alone is insufficient.

1. **Relational Database (PostgreSQL):** This will be our primary data store for all structured, transactional data. It will manage users, workflow definitions, execution states, audit logs, and HITL tasks. PostgreSQL is the ideal choice due to its robustness, transactional integrity (ACID compliance), and powerful features like JSONB support, which we will leverage heavily. Prisma ORM will be used to interact with this database.
2. **Vector Database (Qdrant):** This specialized database is required for the semantic search and Retrieval-Augmented Generation (RAG) capabilities [projects.ai\_search]. It is designed for the sole purpose of storing high-dimensional vector embeddings and performing ultra-fast similarity searches with advanced filtering. Storing and querying vector embeddings in Postgres (via extensions like `pgvector`) is possible, but a dedicated solution like Qdrant provides superior performance, scalability, and features tailored for our needs.

## 1. Relational Database Schema (PostgreSQL with Prisma)

This schema translates the core business logic into a set of interconnected models. The structure is designed to be multi-tenant, secure, and fully auditable.

```
// This is your Prisma schema file,
// learn more about it in the docs: https://pris.ly/d/prisma-schema

generator client {
  provider = "prisma-client-js"
}

datasource db {
  provider = "postgresql"
  url      = env("DATABASE_URL")
}

// =====
```

```

// Core Tenancy & User Management
// =====

/// Represents an enterprise client's isolated environment. All data is scoped to a Tenant
model Tenant {
    id      String @id @default(cuid())
    name    String @unique
    users   User[]

    workflowDefinitions WorkflowDefinition[]
    workflowRuns         WorkflowRun[]
    credentials          Credential[]
    documents            Document[]

    createdAt DateTime @default(now())
    updatedAt DateTime @updatedAt
}

/// Represents a user within a Tenant.
model User {
    id      String @id @default(cuid())
    email   String @unique
    name    String?
    password String // Will be a hashed password

    role Role @default(OPERATOR)

    tenantId String
    tenant    Tenant @relation(fields: [tenantId], references: [id])

    assignedHITLTasks HITLTask[]
    auditLogs         AuditLog[]

    createdAt DateTime @default(now())
    updatedAt DateTime @updatedAt
}

enum Role {
    ADMINISTRATOR /// Manages users, billing, and system settings
    DESIGNER      /// Creates and edits workflows
    OPERATOR      /// Runs workflows and handles HITL tasks
}

// =====
// Workflow Definition & Composition
// =====

/// Stores the design and structure of a workflow from the Workflow Studio.
model WorkflowDefinition {
    id      String @id @default(cuid())
    name    String
    description String?
    status  WorkflowStatus @default(DRAFT)

    /// The JSON representation of the workflow graph (nodes, edges, agent configs).

```

```

definitionGraph Json

tenantId String
tenant    Tenant @relation(fields: [tenantId], references: [id])

runs WorkflowRun[]

createdAt DateTime @default(now())
updatedAt DateTime @updatedAt
}

enum WorkflowStatus {
    DRAFT
    ACTIVE
    ARCHIVED
}

// =====
// Workflow Execution & State
// =====

/// An individual, live instance of a running WorkflowDefinition.
model WorkflowRun {
    id String @id @default(cuid())

    workflowDefinitionId String
    workflowDefinition    WorkflowDefinition @relation(fields: [workflowDefinitionId], refer

    tenantId String
    tenant    Tenant @relation(fields: [tenantId], references: [id])

    status RunStatus @default(PENDING)

    /// The current data payload being passed through the workflow.
    payload Json?
    /// The internal state of the workflow engine for this run (e.g., current step).
    state    Json?

    startedAt    DateTime?
    completedAt  DateTime?

    hitlTasks HITLTask[]
    auditLogs  AuditLog[]

    createdAt DateTime @default(now())
    updatedAt  DateTime @updatedAt
}

enum RunStatus {
    PENDING
    RUNNING
    COMPLETED
    FAILED
    PAUSED_FOR_APPROVAL
    PAUSED_WITH_ERROR

```

```

}

// =====
// Human-in-the-Loop (HITL) & Feedback
// =====

/// Represents a task that requires human review and action.
model HITLTask {
    id String @id @default(cuid())

    workflowRunId String
    workflowRun WorkflowRun @relation(fields: [workflowRunId], references: [id])

    assignedToId String
    assignedTo User @relation(fields: [assignedToId], references: [id])

    status HITLStatus @default(PENDING)

    /// The data that needs to be reviewed by the human.
    taskData Json

    /// The original agent output, logged for feedback analysis.
    originalAgentOutput Json?

    /// The corrected data from the human, if edited. This is critical for the learning loop.
    humanCorrectedOutput Json?

    rejectionReason String?

    createdAt DateTime @default(now())
    resolvedAt DateTime?
}

enum HITLStatus {
    PENDING
    APPROVED
    REJECTED
    EDITED_AND_APPROVED
}

// =====
// Auditing & Logging
// =====

/// An immutable log of every significant action in the system.
model AuditLog {
    id String @id @default(cuid())

    workflowRunId String
    workflowRun WorkflowRun @relation(fields: [workflowRunId], references: [id])

    /// Can be a User or null if the action was taken by the system/agent.
    actorId String?
    actor User? @relation(fields: [actorId], references: [id])
}

```

```

actorType String @default("SYSTEM") // e.g., "USER", "AGENT", "SYSTEM"

action      String // e.g., "WORKFLOW_STARTED", "STEP_APPROVED", "AGENT_ERROR"
details     Json?  // Contextual details about the action

timestamp   DateTime @default(now())
}

// =====
// Integration & Document Management
// =====

/// Stores encrypted credentials for third-party systems.
model Credential {
    id          String @id @default(cuid())
    tenantId    String
    tenant      Tenant @relation(fields: [tenantId], references: [id])

    type        String // e.g., "OAUTH2", "API_KEY"
    serviceName String // e.g., "SALESFORCE", "GMAIL"

    /// The credential value, encrypted at the application layer before being stored.
    encryptedValue String

    createdAt   DateTime @default(now())
}

/// Represents a document indexed in the vector DB for RAG.
/// This model provides a relational link to the vectorized data.
model Document {
    id String @id @default(cuid())

    tenantId String
    tenant    Tenant @relation(fields: [tenantId], references: [id])

    sourceUri  String @unique // e.g., sharepoint://... or file://...
    fileName   String
    fileType   String // e.g., "PDF", "DOCX"

    status      IndexingStatus @default(PENDING)
    lastIndexed DateTime?

    createdAt   DateTime @default(now())
}

enum IndexingStatus {
    PENDING
    INDEXING
    COMPLETED
    FAILED
}

```

## 2. Vector Database "Schema" (Qdrant)

A vector database doesn't have a rigid schema in the same way a relational database does. Instead, we define the structure of the data points we will store in a "collection." For our RAG and semantic search feature, each data point will represent a chunk of a document.

Here is the logical structure for a data point stored in our Qdrant collection:

```
{
  // Unique identifier for the vector point (e.g., a UUID).
  "id": "a1b2c3d4-e5f6-7890-1234-567890abcdef",

  // The high-dimensional vector embedding of the 'content' field.
  // This is what Qdrant uses for similarity search.
  "vector": [0.0123, -0.0456, ..., 0.0789],

  // The metadata payload. This is critical for filtering and security.
  "payload": {
    // The original text chunk from the document.
    "content": "The Q3 2025 sales target for the West region is $2.5 million, an increase",

    // Link back to the relational database for more context.
    "document_id": "doc_2alkj34hkljha",

    // The unique ID of the tenant this data belongs to.
    // **This is a mandatory filter for ALL queries to ensure data isolation.**
    "tenant_id": "tnt_1a2b3c4d5e6f",

    // The source of the document chunk.
    "source_uri": "sharepoint://FinanceDocs/2025/Q3_Sales_Targets.docx",

    // **Security-critical field for permission-aware retrieval.**
    // An array of roles or groups that are allowed to see this data chunk.
    "allowed_roles": ["Finance", "SalesManager", "Executive"],

    // Other useful metadata for filtering.
    "author": "john.doe@enterprise.com",
    "created_at_source": "2025-07-15T10:00:00Z"
  }
}
```

This hybrid schema provides a robust, scalable, and secure foundation to build the full suite of features defined in the business logic for the Enterprise Agentic Workflow Engine.

✱